

# Blockchain on vSphere

By VMware

## Introduction

Blockchain is an emerging technology which has been gaining traction globally during the past few years. Industries like finance, logistics, IoT, are actively working on research and pilot projects using blockchain.

Despite the attention to blockchain, the installation and management of a blockchain service is very complicated and requires sophisticated domain knowledge. Blockchain on vSphere (BoV) is a tool for users to deploy Hyperledger Fabric v1.0 on vSphere platform. With only a few commands, a cluster of Fabric node is up and running on vSphere. The blockchain developers can quickly focus on implementation of the business logic.

Hyperledger is an open source project hosted by Linux Foundation. It was created to advance cross-industry blockchain technologies. It is by far the most mature open source project for enterprise to try and use distributed ledger technology(DLT).

## Use Case

There are three personas of a blockchain service: Cloud Admin, Blockchain Admin, Blockchain Developer. They collaborate at different levels of the system. The Cloud Admin provisions and monitor the infrastructure such as vSphere. The Blockchain Admin manages the blockchain platform (Hyperledger Fabric). The Blockchain Developer focuses on application development by using the blockchain platform.

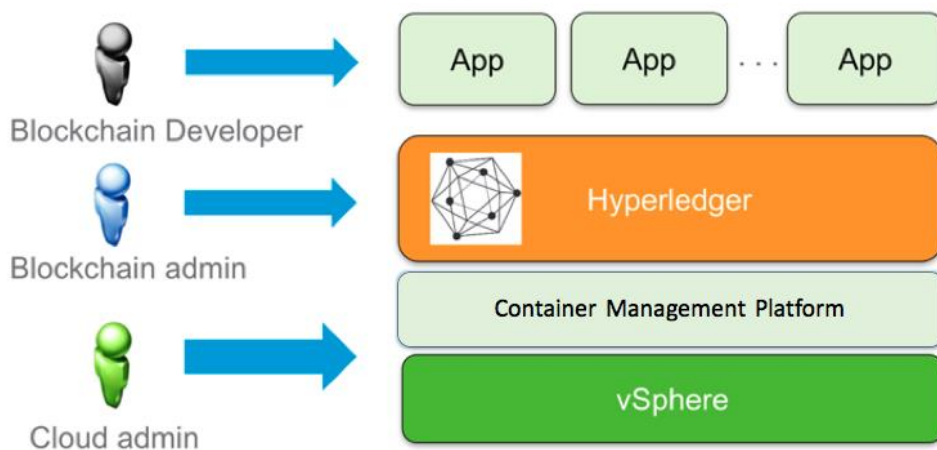


Figure 1

## Architecture

Hyperledger Fabric is a distributed system implemented using containers. It can be deployed on the platform that support OCI container standard. Kubernetes will be used here to manage the Fabric containers. We have used the below architecture for Fabric v1.0:

- Use namespaces to maintain different Fabric organizations.
- Customizable number of peers in an organization.
- Isolation through Persistent Volume.

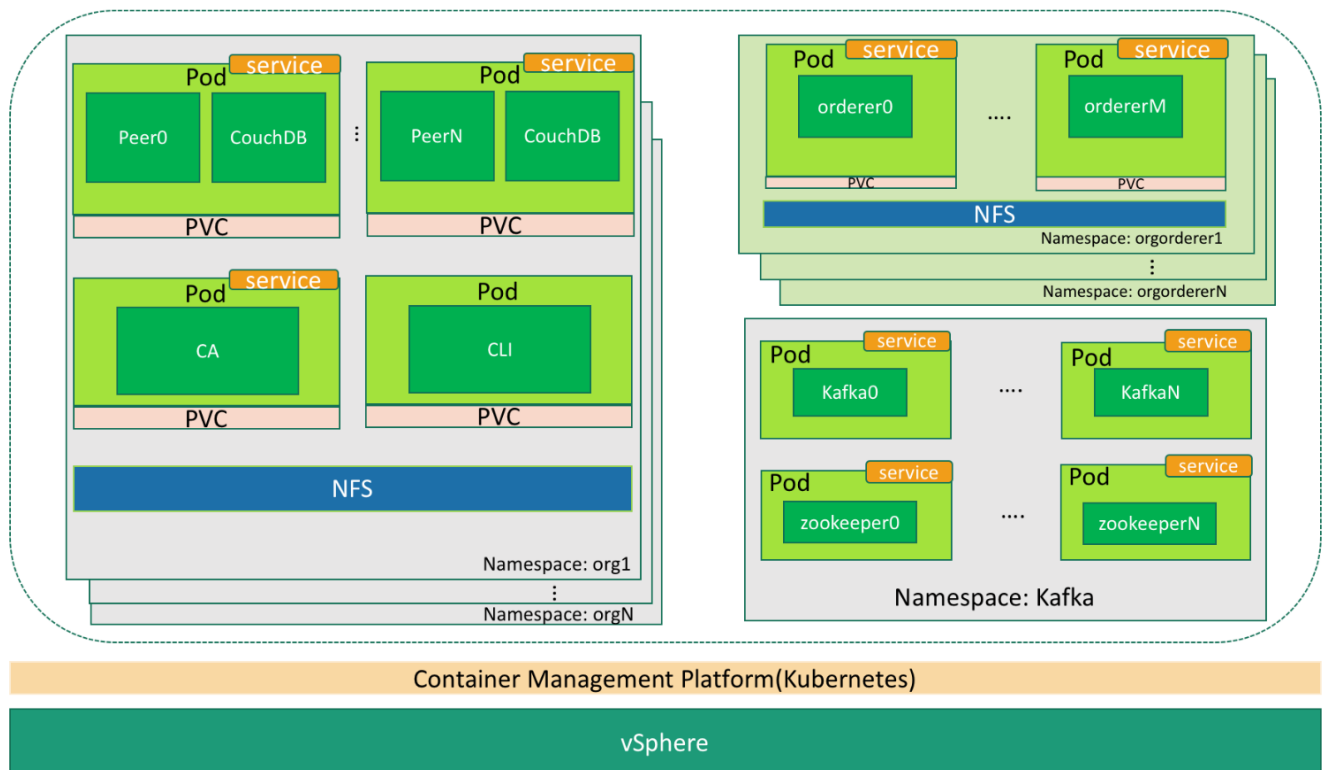


Figure 2

## Deployment Instructions

### Prerequisites:

- 1) vCenter 6.0+, with at least one ESXi host.
- 2) A NFS server for storing configuration files of the Fabric cluster.
- 3) Internet connection is required during installation.
- 4) A Linux host with Docker 1.11+ installed to run commands.
- 5) Python3.5 installed on the Linux host (with [PyYAML](#) module).

After you downloaded the package of BoV, you can follow the below steps to install Fabric 1.0 on vSphere. If you choose to use an existing Kubernetes cluster, you can start from Step 3.

### 1. Prepare vSphere environment

In vCenter, configure or create the below resources for Kubernetes:

- ❖ Create a resource pool for the Kubernetes, such as kub-pool;
- ❖ Select a datastore used by the Kubernetes, such as datastore1;
- ❖ Select a network for the Kubernetes nodes such as VM Network. The network must have DHCP service to provide dynamic IP address and can connect to internet.

## 2. Deployment Kubernetes

We will deploy Kubernetes by using open source project *Kubernetes Anywhere* (<https://github.com/kubernetes/kubernetes-anywhere>).

- ❖ Download this [OVA](https://github.com/kubernetes/kubernetes-anywhere/blob/master/phase1/vsphere/README.md) template file and import it into vCenter. You will get a VM template named "KubernetesAnywhereTemplatePhotonOS". More information about importing the OVA can be found here: <https://github.com/kubernetes/kubernetes-anywhere/blob/master/phase1/vsphere/README.md>

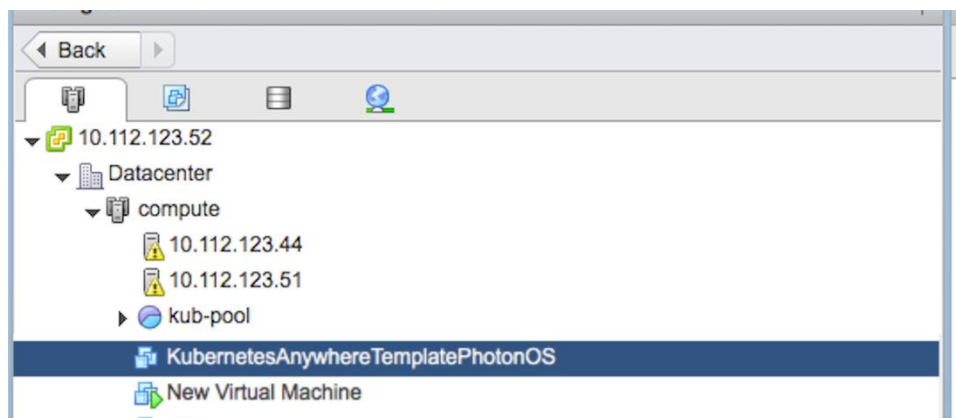


Figure 3

- ❖ On the Linux host, run the deployment container of *Kubernetes Anywhere*, you will stay inside the container's command prompt after the commands:

```
$ docker pull cnastorage/kubernetes-anywhere
```

```
$ docker run -it --rm --env="PS1=[container]:\w> " --net=host cnastorage/kubernetes-anywhere:v1 /bin/bash
```

- ❖ Inside the container start to deploy Kubernetes:
  - [container]:/opt/kubernetes-anywhere> make deploy
  - Input parameters according to the prompts. The following is an example:( to accept the default value, just press the enter key.)
    - number of nodes (phase1.num\_nodes) [4] (NEW) 4
    - cluster name (phase1.cluster\_name) [kubernetes] (NEW) **kubernetes**
    - cloud provider: gce, azure or vsphere (phase1.cloud\_provider) [gce] (NEW) **vsphere**
    - vCenter URL Ex: 10.192.10.30 or myvcenter.io (phase1.vSphere.url) [] (NEW) **10.112.123.52**
    - vCenter port (phase1.vSphere.port) [443] (NEW) **443**
    - vCenter username (phase1.vSphere.username) [] (NEW) **administrator@vsphere.local**
    - vCenter password (phase1.vSphere.password) [] (NEW) **MyPassword#3**
    - Does host use self-signed cert (phase1.vSphere.insecure) [Y/n/?] (NEW) **Y**

Please note only input Y when your VC use a self-singed certificate(Please aware that use self-signed certificate are prone to the man in the middle attack).

- Datacenter (phase1.vSphere.datacenter) [datacenter] (NEW) **Datacenter**
- Datastore (phase1.vSphere.datastore) [datastore] (NEW) **datastore1**

- Deploy Kubernetes Cluster on 'host' or 'cluster' (phase1.vSphere.placement) [cluster] (NEW) **cluster**
- vspherecluster (phase1.vSphere.cluster) [] (NEW) **compute**
- Do you want to use the resource pool created on the host or cluster? [yes, no] (phase1.vSphere.userresourcepool) [no] (NEW) **yes**
- Name of the Resource Pool. If Resource pool is enclosed within another Resource pool, specify pool hierarchy as ParentResourcePool/ChildResourcePool (phase1.vSphere.resourcepool) (NEW) **kub-pool**
- Number of vCPUs for each VM (phase1.vSphere.vcpu) [1] (NEW) **4**
- Memory for each VM (phase1.vSphere.memory) [2048] (NEW) **4096**
- Network for each VM (phase1.vSphere.network) [VM Network] (NEW) **VM Network**
- Name of the template VM imported from OVA. If Template file is not available at the destination location specify vm path (phase1.vSphere.template) [KubernetesAnywhereTemplatePhotonOS.ova] (NEW) **KubernetesAnywhereTemplatePhotonOS**
- Flannel Network (phase1.vSphere.flannel\_net) [172.1.0.0/16] (NEW)
- **\*\* Phase 2: Node Bootstrapping\***
- installer container (phase2.installer\_container) [docker.io/cnastorage/k8s-ignition:v2] (NEW)
- docker registry (phase2.docker\_registry) [gcr.io/google-containers] (NEW)
- kubernetes version (phase2.kubernetes\_version) [v1.6.5] (NEW)
- bootstrap provider (phase2.provider) [ignition] (NEW)
- **\*\* Phase 3: Deploying Addons.\***
- Run the addon manager? (phase3.run\_addons) [Y/n/?] (NEW)
- Run kube-proxy? (phase3.kube\_proxy) [Y/n/?] (NEW)
- Run the dashboard? (phase3.dashboard) [Y/n/?] (NEW)
- Run heapster? (phase3.heapster) [Y/n/?] (NEW)
- Run kube-dns? (phase3.kube\_dns) [Y/n/?] (NEW)
- Run weave-net? (phase3.weave\_net) [N/y/?] (NEW) **N**

- Wait for the Kubernetes cluster to be created. Use the following two commands to check Kubernetes cluster status:

```
$ export KUBECONFIG=phase1/vsphere/.tmp/kubeconfig.json
$ kubectl cluster-info
```

It will display the cluster information similar to the below:

```
Kubernetes master is running at https://10.112.122.30
Heapster is running at https://10.112.122.30/api/v1/proxy/namespaces/kube-system/services/heapster
KubeDNS is running at https://10.112.122.30/api/v1/proxy/namespaces/kube-system/services/kube-dns
```

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.

Figure 4

- Copy the content of the file **phase1/vsphere/.tmp/kubeconfig.json** to your **Linux host** and save the content to the **~/ .kube/config** . You may need to create **.kube** directory if it does not exist on the Linux host:

```
$ scp phase1/vsphere/.tmp/kubeconfig.json username@LinuxHost:~/ .kube/config
```

- Exit the container and return to the Linux host.

- On your Linux host run the following command:

```
$ kubectl cluster-info
```

You will get the same result as Figure 4.

- ❖ Configure the DNS used by Docker on all Kubernetes worker nodes.

The Fabric creates a Docker container to run chaincode which is out of the Kubernetes' control. Therefore, Docker daemon needs to include the correct DNS information of the Kubernetes network. It should include both the Kubernetes' DNS and the worker's DNS that used to reach the internet.

Make the follow change on all your Kubernetes worker nodes. (In this example, it includes node1, node2, node3 and node4)

1) Edit the `/etc/default/docker` file and change its content similar to the following:

```
DOCKER_OPTS="--bip=172.1.42.1/24 --ip-masq=false --mtu=1472 --dns=10.0.0.10 --  
dns=10.117.0.1 --dns-search default.svc.cluster.local --dns-search svc.cluster.local \  
--dns-opt ndots:2 --dns-opt timeout:2 --dns-opt attempts:2 "
```

NOTE: In the above example, Kubernetes Anywhere set **10.0.0.10** as the DNS server of Kubernetes network. **10.117.0.1** is the DNS server of the worker node's network. Please replace the value based on your network configuration.

If you have a proxy server for Docker to pull images, please also add them in the `/etc/default/docker` file, e.g. :

```
HTTP_PROXY=http://yourproxyserver:3128  
HTTPS_PROXY=https://yourproxyserver:3128
```

2) Restart the Docker service to allow the changes to take effect:

```
$ systemctl dockerd-daemon restart
```

```
$ systemctl dockerd restart
```

Now the Kubernetes installation has been completed.

❖ Update software in Kubernetes worker nodes.

As the kubernetes worker VMs are base on Photon ova which is created on Mar/2017 your are suggested to run `"tdnf makecache && tdnf distro-sync"` to keep the software keep to latested.

The default login info for the worker node is: root/kubernetes and your are strongly recommend to change those default password to a more secure one.

### 3. Deploy blockchain platform (Fabric)

❖ You need to setup a NFS service and export a shared directory (e.g. `/opt/share`). You can check the setting on your NFS server (10.112.122.9 in this example)

```
root@ubuntu:~# showmount -e 10.112.122.9
```

```
Export list for 10.112.122.9:
```

`/opt/share *`

The NFS client needs to have read/write access to the `/opt/share` folder. If there is no authentication required by NFS (i.e. anonymous access), the folder's owner and group needs to be changed to **nobody:nogroup**. Otherwise the Kubernetes pods will encounter permission error. You can simply run the below command on the NFS server:

```
$ chmod -R nobody:nogroup /opt/share
```

- ❖ Mount the `/opt/share` to your Linux host:

```
$ mount 10.112.122.9:/opt/share /opt/share
```

- ❖ Download the BoV package file `Baas.tar` and extract the files.
- ❖ Change the current directory to `./baas` and run the command to download tools required by Fabric. Two tools, `cryptogen` and `configtxgen`, will be saved in `./bin` directory.:

```
$ cd ./baas
```

```
$ curl
```

```
https://nexus.hyperledger.org/content/repositories/releases/org/hyperledger/fabric/hyperledger-fabric/linux-amd64-1.0.0/hyperledger-fabric-linux-amd64-1.0.0.tar.gz | tar xz
```

```
$ chmod +x ./bin/*
```

- ❖ In the `setupCluster/templates/` directory, update two template files with your NFS server's IP address.

1) `fabric_1_0_template_pod_cli.yaml`

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: $artifactsName
spec:
  capacity:
    storage: 500Mi
  accessModes:
    - ReadWriteMany
  nfs:
    path: /opt/share/channel-artifacts
    server: 10.112.122.9 # change to you
---
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  namespace: $namespace
  name: $artifactsName
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 10Mi
```

Figure 5

2) `fabric_1_0_template_pod_namespace.yaml`

```

apiVersion: v1
kind: Namespace
metadata:
  name: $org
---
apiVersion: v1
kind: PersistentVolume
metadata:
  name: $pvName
spec:
  capacity:
    storage: 500Mi
  accessModes:
    - ReadWriteMany
  nfs:
    path: $path
    server: 10.112.122.9 #ch

```

Figure 6

- ❖ The file `setupCluster/cluster-config.yaml` contains the definition of the blockchain service topology. You can modify it based on your requirement. Below is an example:

```

Tenant: f-1

OrdererOrgs:
- Name: Orderer
  Domain: orgorderer
  Template:
    Count: 1          # 1 orderer, change this number to create more orderers

PeerOrgs:
- Name: Org1
  Domain: org1        # the domain name of Org1 is org1
  Template:
    Count: 2          # Org1 will have 2 peer nodes.
  Users:
    Count: 1
- Name: Org2
  Domain: org2
  Template:
    Count: 2          # Org2 has 2 peer nodes
  Users:
    Count: 1
- Name: Org3
  Domain: org3
  Template:
    Count: 3          # Org3 has 3 peer nodes
  Users:
    Count: 1

```

- ❖ Change directory to `baas/setupCluster` and run `generateAll.sh` to create all the Fabric configuration files and Kubernetes pods definition files for the blockchain service.

```
$ cd baas/SetupCluster
```

```
$ sudo bash generateAll.sh
```

You can view the generated configuration files at the `/opt/share` directory.

- ❖ Change directory to the baas/setupCluster/transform and run the following command to deploy Fabric as pods on Kubernetes:

```
$ python3.5 run.py
```

- ❖ Check the created blockchain pods:

```
$ kubectl get pods --all-namespaces
```

```
root@master1 ~]# kubectl get pods --all-namespaces -o wide
NAMESPACE   NAME                                     READY   STATUS    RESTARTS   AGE   IP             NODE
kube-system  etcd-server-master                     1/1     Running   1           24d   10.112.122.30  master
kube-system  heapster-v1.2.0-867844254-wghgj       2/2     Running   14          24d   172.1.1.86.3   node3
kube-system  kube-apiserver-master                  1/1     Running   1           24d   10.112.122.30  master
kube-system  kube-controller-manager-master         1/1     Running   1           24d   10.112.122.30  master
kube-system  kube-dns-v19-l0bdz                     3/3     Running   15          18d   172.1.1.86.2   node3
kube-system  kube-dns-v19-nglhc                     3/3     Running   21          24d   172.1.1.12.2   node2
kube-system  kube-proxy-044r8                       1/1     Running   7           24d   10.112.122.32  node2
kube-system  kube-proxy-0786x                       1/1     Running   2           17d   10.112.122.31  node4
kube-system  kube-proxy-dvnt2                       1/1     Running   8           24d   10.112.122.48  node3
kube-system  kube-proxy-n6rvx                       1/1     Running   1           24d   10.112.122.30  master
kube-system  kube-proxy-xrm5q                       1/1     Running   9           24d   10.112.122.40  node1
kube-system  kube-scheduler-master                  1/1     Running   1           24d   10.112.122.30  master
kube-system  kubernetes-dashboard-1019458639-1cx1v  1/1     Running   1           16d   172.1.1.42.2   node4
org1-f-1     ca-3347986348-9jvht                    1/1     Running   0           1h    172.1.73.5     node1
org1-f-1     cli-1569835662-01c5z                   1/1     Running   0           1h    172.1.12.5     node2
org1-f-1     peer0-org1-f-1-1343141255-h8kgk        2/2     Running   0           1h    172.1.86.6     node3
org1-f-1     peer1-org1-f-1-2603922830-35q6c        2/2     Running   0           1h    172.1.12.6     node2
org2-f-1     ca-2708682628-qpz64                    1/1     Running   0           1h    172.1.73.3     node1
org2-f-1     cli-2586364563-vc1mr                   1/1     Running   0           1h    172.1.42.4     node4
org2-f-1     peer0-org2-f-1-3143546256-9prph        2/2     Running   0           1h    172.1.73.4     node1
org2-f-1     peer1-org2-f-1-110343575-06pvc         2/2     Running   0           1h    172.1.42.5     node4
org3-f-1     ca-349255610-628k1                     1/1     Running   0           1h    172.1.12.3     node2
org3-f-1     cli-3602893464-7f6g1                   1/1     Running   0           1h    172.1.73.2     node1
org3-f-1     peer0-org3-f-1-649967001-0v813         2/2     Running   0           1h    172.1.12.4     node2
org3-f-1     peer1-org3-f-1-1910748576-1jlbv        2/2     Running   0           1h    172.1.86.5     node3
org3-f-1     peer2-org3-f-1-3171530151-6whd0        2/2     Running   0           1h    172.1.42.3     node4
orgorderer-f-1  orderer-orgorderer-f-1-73543963-plc1z  1/1     Running   0           1h    172.1.86.4     node3
```

Figure 7

You can also check it through the Kubernetes dashboard UI (see Figure 8). If you don't know how to access the UI, just follow these steps:

```
# Get NodePort mapping
```

```
kubectl describe service kubernetes-dashboard --namespace=kube-system | grep -i NodePort
```

Output:

```
      Type:      NodePort
NodePort: <unset> 31165/TCP
```

```
# Get node it is running on
```

```
kubectl get pods --namespace=kube-system | grep -i dashboard
```

Output:

```
kubernetes-dashboard-1019458639-1cx1v 1/1 Running 1
```

```
kubectl describe pod kubernetes-dashboard-1763797262-fzla9 --namespace=kube-system | grep Node
```

Output:

```
Node:      node4/10.112.122.31
Node-Selectors: <none>
```

```
# Select the public IP for the node via or use govc or vCenter UI
```

```
kubectl describe node node1 | grep Address
```

```
# Open the <IP Addr>:<NodePort> in a browser
```



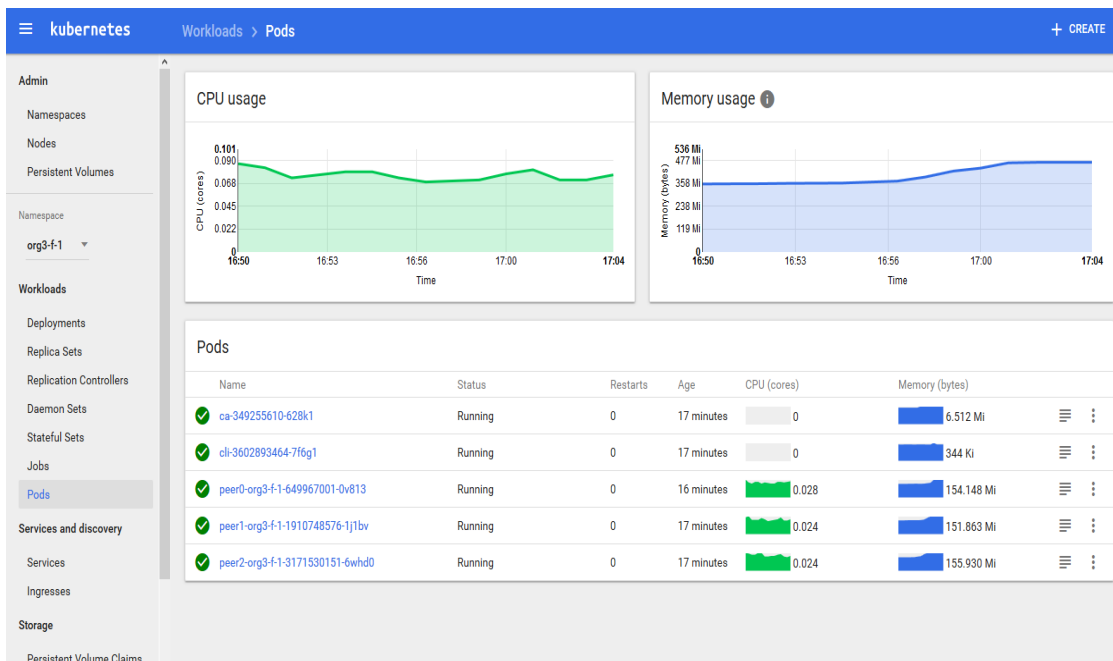


Figure 8

#### 4. Verify the environment by running a sample chaincode

Please follow the next few steps to create channel artifacts that are shared by all pods (two organizations in the example):

- Change working path to `baas/setupCluster/` :  

```
$ cd baas/setupCluster/
```
- Create mychannel by using `configtxgen`:  

```
$ ../bin/configtxgen -profile TwoOrgsChannel -outputCreateChannelTx ./channel-artifacts/channel.tx -channelID mychannel
```
- Create configuration transaction for updating peer0.org1 to be the anchor peer of mychannel:  

```
$ ../bin/configtxgen -profile TwoOrgsChannel -outputAnchorPeersUpdate ./channel-artifacts/Org1MSPanchors.tx -channelID mychannel -asOrg Org1MSP
```
- Create configuration transaction for updating peer0.org2 to be anchor peer of mychannel:  

```
$ ../bin/configtxgen -profile TwoOrgsChannel -outputAnchorPeersUpdate ./channel-artifacts/Org2MSPanchors.tx -channelID mychannel -asOrg Org2MSP
```

- Copy `./channel-artifacts` directory into `/opt/share/` so that all cli pods can see it:

```
$ cp -r ./channel-artifacts /opt/share
```

Download the chaincode example from Hyperledger [Fabric project](#) . Then copy the `chaincode_example02` directory into `/opt/share/channel-artifacts/` .

In the below commands, it first creates a channel, then it joins peer0 of Org1 into the channel. Next, it installs and instantiates the chaincode with two key/value pairs: `a=100` and `b=200`. After that, it adds peer0 of Org2 to the channel and installs the chaincode. Next it creates a transaction to transfer value 10 from `a` to `b`. It then queries the value of `a` which should be displayed as 90.

❖ Command for Org1:

1. Find out the cli pod instance of Org1:

```
$ kubectl get pods --namespace=org1-f-1
root@master [ ~ ]# kubectl get pods --namespace org1-f-1
NAME                                READY    STATUS    RESTARTS   AGE
ca-3347986348-9jvht                 1/1      Running   0           2h
cli-1569835662-01c5z                1/1      Running   0           2h
peer0-org1-f-1-1343141255-h8kgk     2/2      Running   0           2h
peer1-org1-f-1-2603922830-35q6c     2/2      Running   0           2h
root@master [ ~ ]#
```

Figure 9

As Figure 9 shows, cli of org1 was named `cli-1569835662-01c5z`

2. Enter the cli pod of org1 for a Fabric command prompt:

```
$ kubectl exec -it cli-1569835662-01c5z bash --namespace=org1-f-1
```

3. In the cli pod, create a channel named `mychannel`:

```
$ peer channel create -o orderer0.org:orderer-f-1:7050 -c mychannel -f ./channel-artifacts/channel.tx
```

4. After the channel is created, a file `mychannel.block` from orderer is returned. It will be used by other peers to join the channel later. Figure 10 shows the output of channel creation command:

```
2017-08-07 08:52:11.803 UTC [main] main -> INFO 024 Exiting.....
root@cli-1569835662-01c5z:/opt/gopath/src/github.com/hyperledger/fabric/peer# ls
channel-artifacts  mychannel.block
root@cli-1569835662-01c5z:/opt/gopath/src/github.com/hyperledger/fabric/peer#
```

Figure 10

5. copy `mychannel.block` to NFS shared folder so that other cli can use it:

```
$ cp mychannel.block ./channel-artifacts
```

6. Peer0 of Org1 joins the *mychannel*:

```
$ peer channel join -b ./channel-artifacts/mychannel.block
```

If everything worked fine, the output should look like Figure 11:

```
2017-08-07 09:17:12.475 UTC [channelCmd] executeJoin -> INFO 009 Peer joined the channel!
2017-08-07 09:17:12.476 UTC [main] main -> INFO 00a Exiting.....
```

Figure 11

7. Update peer0 of Org1 as the anchor peer of *mychannel*:

```
$ peer channel update -o orderer0.orgorderer-f-1:7050 -c mychannel -f ./channel-artifacts/Org1MSPanchors.tx
```

8. Install chaincode *mycc* on peer0:

```
$ peer chaincode install -n mycc -v 1.0 -p github.com/hyperledger/fabric/peer/channel-artifacts/chaincode_example02
```

This output indicates the chaincode was installed successfully:

```
2017-08-07 09:23:50.336 UTC [chaincodeCmd] install -> DEBU 00f Installed remotely response:<status:200 payload:"OK" >
2017-08-07 09:23:50.336 UTC [main] main -> INFO 010 Exiting.....
```

Figure 12

9. Instantiate the chaincode *mycc*:

```
$ peer chaincode instantiate -o orderer0.orgorderer-f-1:7050 -C mychannel -n mycc -v 1.0 -c '{"Args":["init","a","100","b","200"]}' -P "OR ('Org1MSP.member','Org2MSP.member')"
```

A chaincode container was created. However, the container is not managed by Kubernetes due to the design of Fabric. The command above use peer0.org1(which is scheduled to node3) to instantiate chaincode. The chaincode container can be found by running "docker ps" at node3 like Figure 13:

```
root@node3 [ ~ ]# docker ps
CONTAINER ID        IMAGE
5acb0198550a        dev-peer0.org1-f-1-mycc-1.0
70fc8398f3d6        e01c2b645f11
kgk_org1-f-1_148b1585-7b4d-11e7-8aa2-0050568fe649_0
f3ca57c3a28c        31bbbec3d853
1-f-1_148b1585-7b4d-11e7-8aa2-0050568fe649_0
068c0784d67e        gcr.io/google_containers/pause-amd64:3.0
```

Figure 13

10. Query the chaincode:

```
$ peer chaincode query -C mychannel -n mycc -c '{"Args":["query","a"]}'
```

Check the chaincode state, the value of "a" should be 100 as the instantiation defined:

```
Query Result: 100
2017-08-07 10:02:03.710 UTC [main] main -> INFO 008 Exiting.....
```

Figure 14

❖ Similar commands can be applied in Org2:

1. Find out the cli pod instance of Org2:

```
$ kubectl get pods --namespace=org2-f-1
```

```
root@master [ ~ ]# kubectl get pods --namespace org2-f-1
NAME                                READY    STATUS    RESTARTS   AGE
ca-2708682628-qpz64                1/1      Running   0           2h
cli-2586364563-vclmr               1/1      Running   0           2h
peer0-org2-f-1-3143546256-9prph    2/2      Running   0           2h
peer1-org2-f-1-110343575-06pvc     2/2      Running   0           2h
```

Figure 15

As Figure 15 showed, cli of org2 was named cli-2586364563-vclmr

2. Enter the cli pod of org2 for a Fabric command prompt:

```
$ kubectl exec -it cli-2586364563-vclmr --namespace=org2-f-1 bash
```

3. In cli pod of org2, Peer0 of org2 joins the channel:

```
$ peer channel join -b ./channel-artifacts/mychannel.block
```

```
2017-08-07 09:17:12.475 UTC [channelCmd] executeJoin -> INFO 009 Peer joined the channel!
2017-08-07 09:17:12.476 UTC [main] main -> INFO 00a Exiting.....
```

Figure 16

4. Update Peer0 of org2 as an anchor peer:

```
$ peer channel update -o orderer0.orgorderer-f-1:7050 -c mychannel -f ./channel-artifacts/Org2MSPanchors.tx
```

5. Install chaincode:

```
$ peer chaincode install -n mycc -v 1.0 -p github.com/hyperledger/fabric/peer/channel-artifacts/chaincode_example02
```

```
2017-08-07 09:23:50.336 UTC [chaincodeCmd] install -> DEBU 00f Installed remotely response:<status:200 payload:"OK" >
2017-08-07 09:23:50.336 UTC [main] main -> INFO 010 Exiting.....
```

Figure 17

6. Query the chaincode:

```
$ peer chaincode query -C mychannel -n mycc -c '{"Args":["query","a"]}'
```

Since chaincode *mycc* was instantiated, this command returns the same "a" value as instantiation defined:

```
Query Result: 100
2017-08-07 10:02:03.710 UTC [main] main -> INFO 008 Exiting....
```

Figure 18

7. Invoke the chaincode to update the ledger of the channel. This creates a transaction to transfer value 10 from *a* to *b*:

```
$ peer chaincode invoke -o orderer0.orgorderer-f-1:7050 -C mychannel -n mycc -c
 '{"Args":["invoke","a","b","10"]}'
```

8. Query "a" again, it should be 90 now:

```
$ peer chaincode query -C mychannel -n mycc -c '{"Args":["query","a"]}'
```

```
Query Result: 90
2017-08-07 10:13:01.232 UTC [main] main -> INFO 008 Exiting....
:~$ cat /dev/null
```

Figure 19

Now you have created a fully functional blockchain service and run a chaincode example on the service.